

Google™





# From Spark Plug to Drive Train: Life of an App Engine Request

Adam Byrtek

Uniwersytet Warszawski, 21 Oct 2010

Based on Google I/O 2009 slides by  
Alon Levi



# Agenda

- **Designing for Scale and Reliability**

- **App Engine: Design Motivations**

- **Life of a Request**

Request for static content

Request for dynamic content

Requests that use APIs

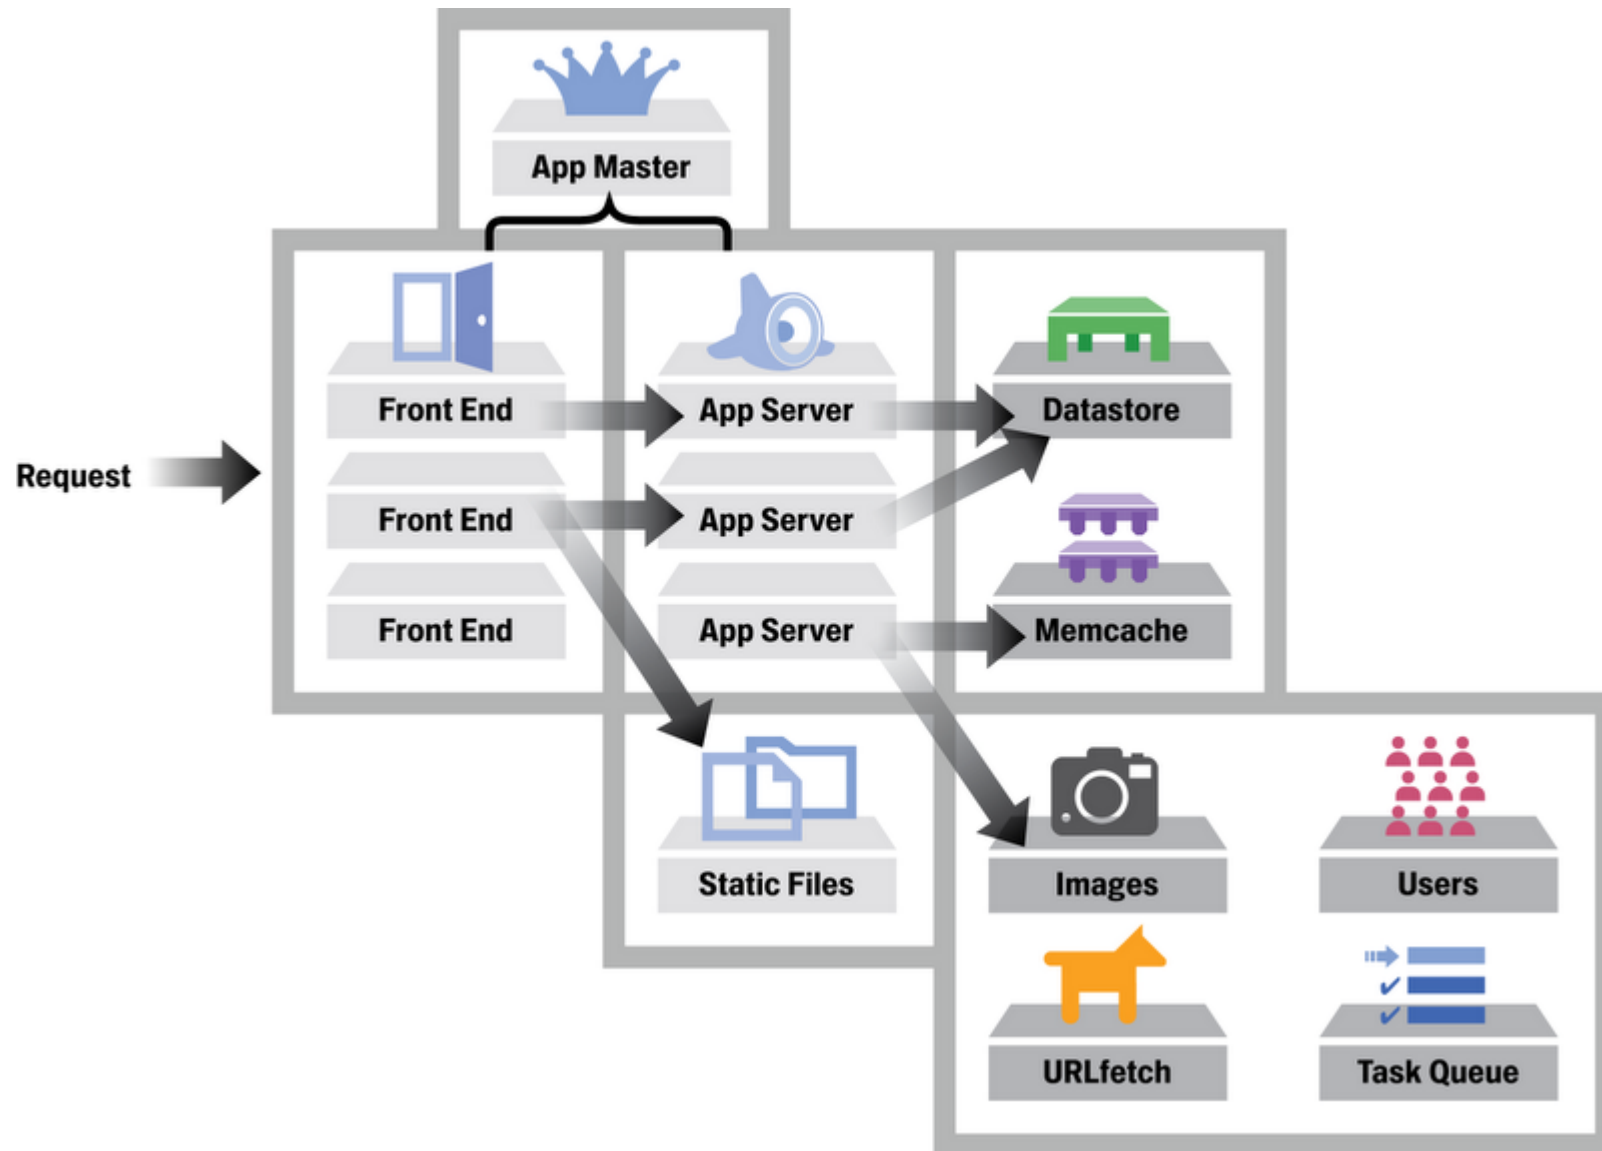
- **App Engine: Design Motivations, Recap**



# Designing for Scale and Reliability




# Google App Engine




# Basic LAMP

- Linux, Apache, MySQL, Programming Language
- **Scalable?**  
Shared machine for database and webserver
- **Reliable?**  
Single point of failure (SPOF)




# Dedicated Database

- Database running on a separate server
- **Requirements**  
Another machine plus additional management
- **Scalable?**  
Up to one web server
- **Reliable?**  
**Two** single points of failure



# Multiple Web Servers




- **Benefits:**  
Grow traffic beyond the capacity of one webserver
- **Requirements:**  
More machines  
Set up load balancing



# Multiple Web Servers


## Load Balancing: DNS Round Robin



- Register list of IPs with DNS
- Statistical load balancing
- DNS record is cached with Time To Live (TTL)
  - TTL may not be respected

# Multiple Web Servers

## Load Balancing: DNS Round Robin




- Register list of IPs with DNS
- Statistical load balancing
- DNS record is cached with Time To Live (TTL)
  - TTL may not be respected

**Now wait for DNS changes to propagate :-)**

# Multiple Web Servers

## Load Balancing: DNS Round Robin




- **Scalable?**

Add more webservers as necessary  
Still I/O bound on one database

- **Reliable?**

Cannot redirect traffic quickly  
Database still SPOF

# Reverse Proxy



- **Benefits:**

Custom Routing

- Specialization
- Application-level load balancing

- **Requirements:**

More machines

Configuration and code for reverse proxies

# Reverse Proxy

## ● Scalable?

Add more web servers

Specialization

Bound by

- Routing capacity of reverse proxy
- One database server

## ● Reliable?

Agile application-level routing

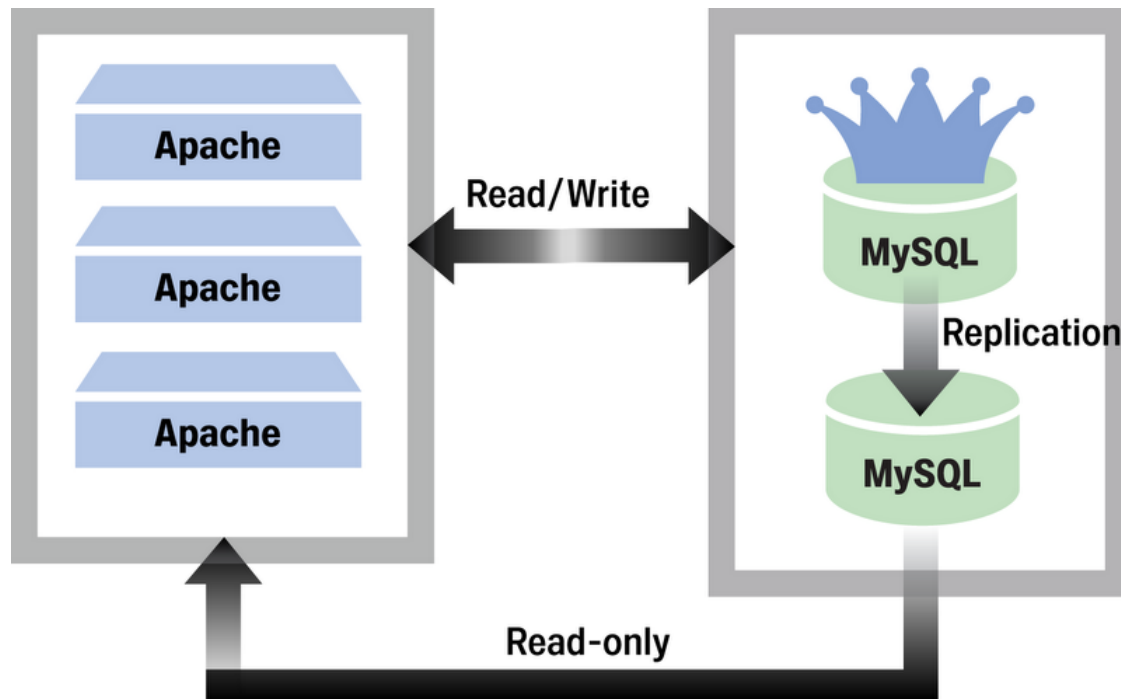
Specialized components are more robust

Multiple reverse proxies requires network-level routing

- DNS Round Robin (again)
- Fancy network routing hardware

Database is still SPOF

# Master-Slave Database



- **Benefits:**

Better read throughput  
Invisible to application

- **Requirements:**

Even more machines  
Changes to MySQL

# Master-Slave Database

- **Scalable?**

Scales read rate with # of servers

- But not writes



What happens eventually?

# Master-Slave Database


- **Reliable?**

Master is SPOF for writes

Master may die before replication



# Partitioned Database




- **Benefits:**

Increase in both read and write throughput

- **Requirements:**

Even more machines  
Lots of management  
Re-architect data model  
Rewrite queries

# The App Engine Stack





# App Engine: Design Motivations



# Design Motivations


- **Build on Existing Google Technology**
- **Provide an Integrated Environment**
- **Encourage Small Per-Request Footprints**
- **Encourage Fast Requests**
- **Maintain Isolation Between Applications**
- **Encourage Statelessness and Specialization**
- **Require Partitioned Data Model**



# Life of a Request: Request for Static Content




# Request for Static Content on Google Network



- Routed to the nearest Google datacenter
- Travels over Google's network
  - Same infrastructure other Google products use
  - Lots of advantages for free

# Request for Static Content

## Routing at the Front End




- **Google App Engine Front Ends**

Load balancing  
Routing

- Frontends route static requests to specialized serving infrastructure

# Request for Static Content

## Static Content Servers



- **Google Static Content Serving**

Built on shared Google Infrastructure

- Static files are physically separate from code files
- How are static files defined?



# Request for Static Content

## What content is static?

### Java Runtime: appengine-web.xml


```
...  
<static>  
  <include path="/**/*.png" />  
  <exclude path="/data/**/*.png" />  
</static>  
...
```

### Python Runtime: app.yaml

```
...  
- url: /images  
static_dir: static/images  
OR  
- url: /images/(.*)  
static_files: static/images/\1  
upload: static/images/(.*)  
...
```

# Request For Static Content

## Response to the user



- Back to the Front End and out to the user
- Specialized infrastructure
  - App runtimes don't serve static content



# Life of a Request: Request for Dynamic Content



# Request for Dynamic Content: New Components

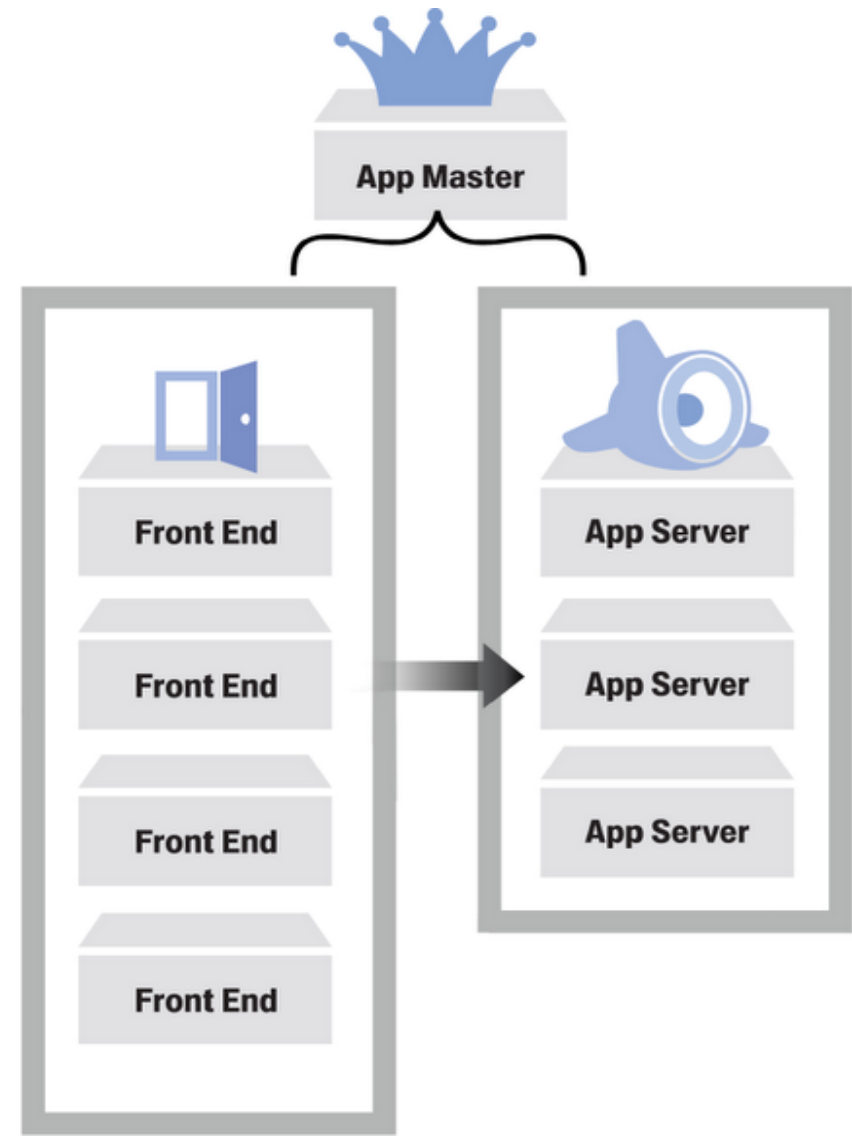
## App Servers and App Master

- **App Servers**

Serve dynamic requests  
Where your code runs

- **App Master**

Schedules applications  
Informs Front Ends



# Request for Dynamic Content: Appservers

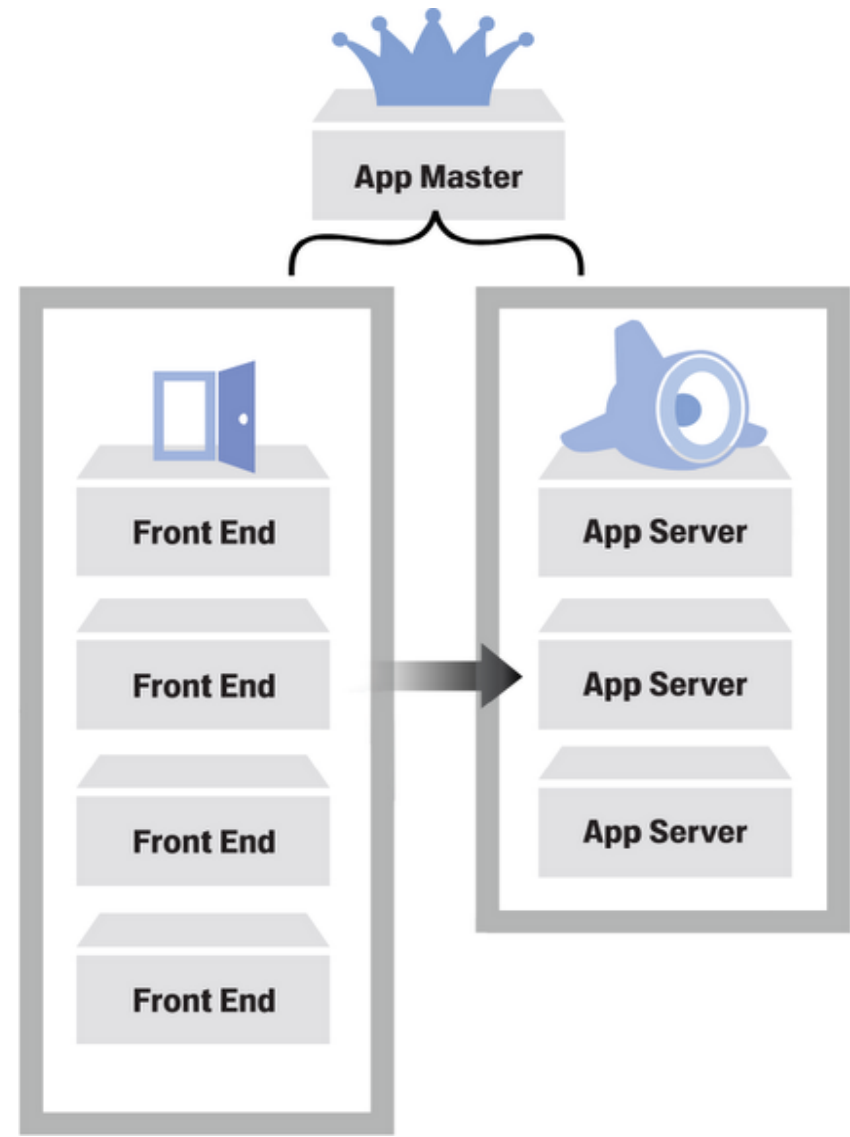
## What do they do?



- Many applications
- Many concurrent requests
  - Smaller app footprint + fast requests = more apps
- Enforce Isolation
  - Keeps apps safe from each other
- Enforce statelessness
  - Allows for scheduling flexibility
- Service API requests


# Request For Dynamic Content Routing at the Frontend

- Front Ends route dynamic requests to App Servers




# Request for Dynamic Content


## App Server



1. Checks for cached runtime
    - If it exists, no initialization
  2. Execute request
  3. Cache the runtime
    - System is designed to maximize caching
- Slow first request, faster subsequent requests
  - Optimistically cache data in your runtime!



# Life of a Request: Requests accessing APIs



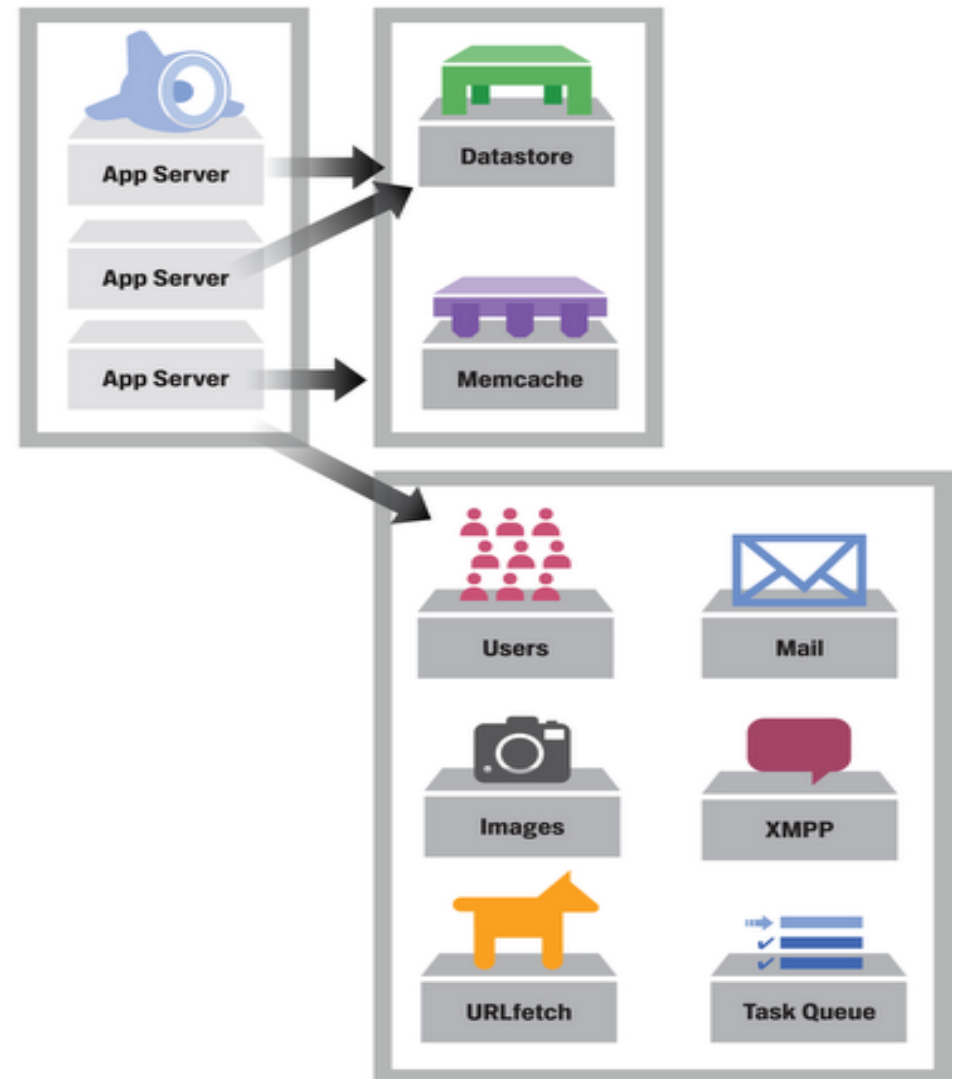


# API Requests




## App Server

1. App issues API call
2. App Server accepts
3. App Server blocks runtime
4. App Server issues call
5. Returns the response

- Use APIs to do things you don't want to do in your runtime, such as...




# APIs



# Memcacheg


**A more persistent in-memory cache**



- Distributed in-memory cache
- memcacheg
  - Also written by Brad Fitzpatrick
  - adds: set\_multi, get\_multi, add\_multi
- Optimistically cache for optimization
- Very stable, robust and specialized

# The App Engine Datastore

## Persistent storage



# The App Engine Datastore

## Persistent storage



- Your data is already partitioned
  - Use Entity Groups
- Explicit Indexes make for fast reads
  - But slower writes
- Replicated and fault tolerant
  - On commit:  $\geq 3$  machines
  - Geographically distributed
- Bonus: Keep globally unique IDs for free


# Other APIs




- GMail



- Gadget API



- Google Accounts



- On roadmap



- Picasaweb



- Google Talk



# App Engine: Design Motivations, Recap



# Provide an Integrated Environment

- **Why?**

Manage all apps together

- **What it means for you:**

Follow best practices

Some restrictions

Use our tools

- **Benefits:**

Use our tools

Admin Console

All of your logs in one place

No machines to configure or manage

Easy deployment



# Encourage Small Per-Request Footprints

- **Why?**

Better utilization of App Servers

Fairness

- **What it means for your app:**

Less Memory Usage

Limited CPU

- **Benefits:**

Better use of resources

# Encourage Fast Requests

- **Why?**

Better utilization of appservers

Fairness between applications

Routing and scheduling agility

- **What it means for your app:**

Runtime caching

Request deadlines

- **Benefits:**

Optimistically share state between requests

Better throughput

Fault tolerance

Better use of resources

# Maintain Isolation Between Apps

- **Why?**

Safety

Predictability

- **What it means for your app:**

Certain system calls unavailable

- **Benefits:**

Security

Performance

# Encourage Statelessness and Specialization

- **Why?**

App Server performance

Load balancing

Fault tolerance

- **What this means for you app:**

Use API calls

- **Benefits:**

Automatic load balancing

Fault tolerance

Less code for you to write

Better use of resources

# Require Partitioned Data Model

- **Why?**

The Datastore

- **What this means for your app:**

Data model + Indexes

Reads are fast, writes are slower

- **Benefits:**

Design your schema once

- No need to re-architect for scalability

More efficient use of cpu and memory



Questions?



Google™

